

Intelligent Multi-Agent Portfolio Optimization System

An AI-Powered Framework for Automated Trading & Risk Management

Ashwini (MA25M006) Dibyendu Sarkar (MA25M011) Jyoti Ranjan Sethi
(MA25M013)

IIT Madras

February 12, 2026

Abstract

We are developing an **intelligent multi-agent system** for automated portfolio optimization that addresses critical limitations in traditional investment management. Our system integrates five specialized applications—risk analysis, real-time monitoring, agent coordination, portfolio optimization, and AI research—working through seven autonomous agents that collaborate to make optimal investment decisions.

Unlike conventional single-model approaches, our multi-agent architecture enables **dynamic adaptation** to market conditions, **transparent decision-making**, and **robust risk management**. Each agent specializes in a specific task while communicating through an orchestration layer that synthesizes diverse inputs into coherent investment strategies.

This 16-week project demonstrates how modern AI techniques can transform portfolio management by combining quantitative models, natural language processing, and real-time data analysis into an interpretable, scalable, and adaptive trading system.

Outline

- 1 Introduction & Motivation
- 2 Project Overview
- 3 System Architecture
- 4 Technical Stack
- 5 Development Timeline
- 6 Expected Outcomes
- 7 Innovation & Research
- 8 Project Deliverables
- 9 Conclusion

The Problem with Traditional Portfolio Management

Key Challenges

- **Static models** cannot adapt to rapidly changing markets
- **Siloed approaches** - risk, optimization, and execution disconnected
- **Limited real-time** monitoring and decision-making
- **Black-box models** lack interpretability
- **Manual intervention** required for strategy adjustments

Our Goal

Build an intelligent, adaptive, and transparent multi-agent system for automated portfolio management

Why Multi-Agent Systems?

Traditional Approach

- Monolithic optimization
- Single decision point
- Limited flexibility
- Hard to debug
- Difficult to update

Multi-Agent Approach

- Specialized agents
- Collaborative decisions
- Highly modular
- Transparent reasoning
- Easy to extend

Multi-agent systems enable better adaptation and interpretability

Our Comprehensive Solution

We are building **FIVE** interconnected applications

- ① **RiskIQ** - Advanced Risk Analysis Engine
- ② **Trading Dashboard** - Real-Time Monitoring Interface
- ③ **Multi-Agent Coordinator** - Agent Orchestration System
- ④ **AI Financial Modeler** - Portfolio Optimization Engine
- ⑤ **AI Research Assistant** - Intelligence & NLP Layer

Integration

These applications work together through a **multi-agent architecture** where seven specialized agents collaborate to make optimal investment decisions

Application 1: RiskIQ - Risk Analysis Engine

What We're Building

Comprehensive real-time risk assessment platform

Key Features:

- Value at Risk (VaR) & Conditional VaR (CVaR) using Monte Carlo
- Dynamic correlation analysis across assets
- Stress testing with customizable scenarios
- Real-time risk alerts and threshold monitoring
- Portfolio sensitivity analysis

Technologies: Python, NumPy, SciPy, QuantLib

Application 2: Trading Dashboard

What We're Building

Interactive web-based monitoring and control interface

Key Features:

- Real-time price charts with technical indicators
- Portfolio performance tracking with attribution
- Position monitoring with live P&L
- Agent activity logs and decision transparency
- Customizable alerts and notifications

Technologies: Streamlit/Gradio, Plotly, React.js

Application 3: Multi-Agent Coordination System

What We're Building

The orchestration layer enabling agent collaboration

Key Features:

- Agent communication protocol with standardized messages
- Consensus-based decision making with weighted voting
- Conflict resolution mechanisms
- Workflow management and task scheduling
- Performance monitoring and agent evaluation

Technologies: LangChain/CrewAI, FastAPI, Redis

Application 4: AI Financial Modeler

What We're Building

Advanced portfolio optimization engine

Key Features:

- Mean-variance optimization (Sharpe ratio maximization)
- Black-Litterman model for market views
- Constraint-based optimization (sectors, ESG, position sizing)
- Dynamic rebalancing with transaction cost optimization
- ML-based return prediction integration

Technologies: PyPortfolioOpt, CVXPY, scikit-learn

Application 5: AI Research Assistant

What We're Building

Intelligent research platform using NLP

Key Features:

- Real-time news sentiment analysis (Transformers)
- Automated company and sector research reports
- Earnings call transcript analysis
- Market regime detection
- Financial question-answering system

Technologies: Hugging Face Transformers, LangChain, OpenAI API

Seven Specialized Agents Working Together

- ① **Market Data Agent** - Real-time data collection & preprocessing
- ② **Research Intelligence Agent** - News & sentiment analysis
- ③ **Alpha Signal Agent** - Trading signal generation
- ④ **Risk Management Agent** - Continuous risk monitoring
- ⑤ **Portfolio Optimization Agent** - Asset allocation optimization
- ⑥ **Execution Agent** - Order placement & position tracking
- ⑦ **Orchestration Agent** - Coordination & consensus

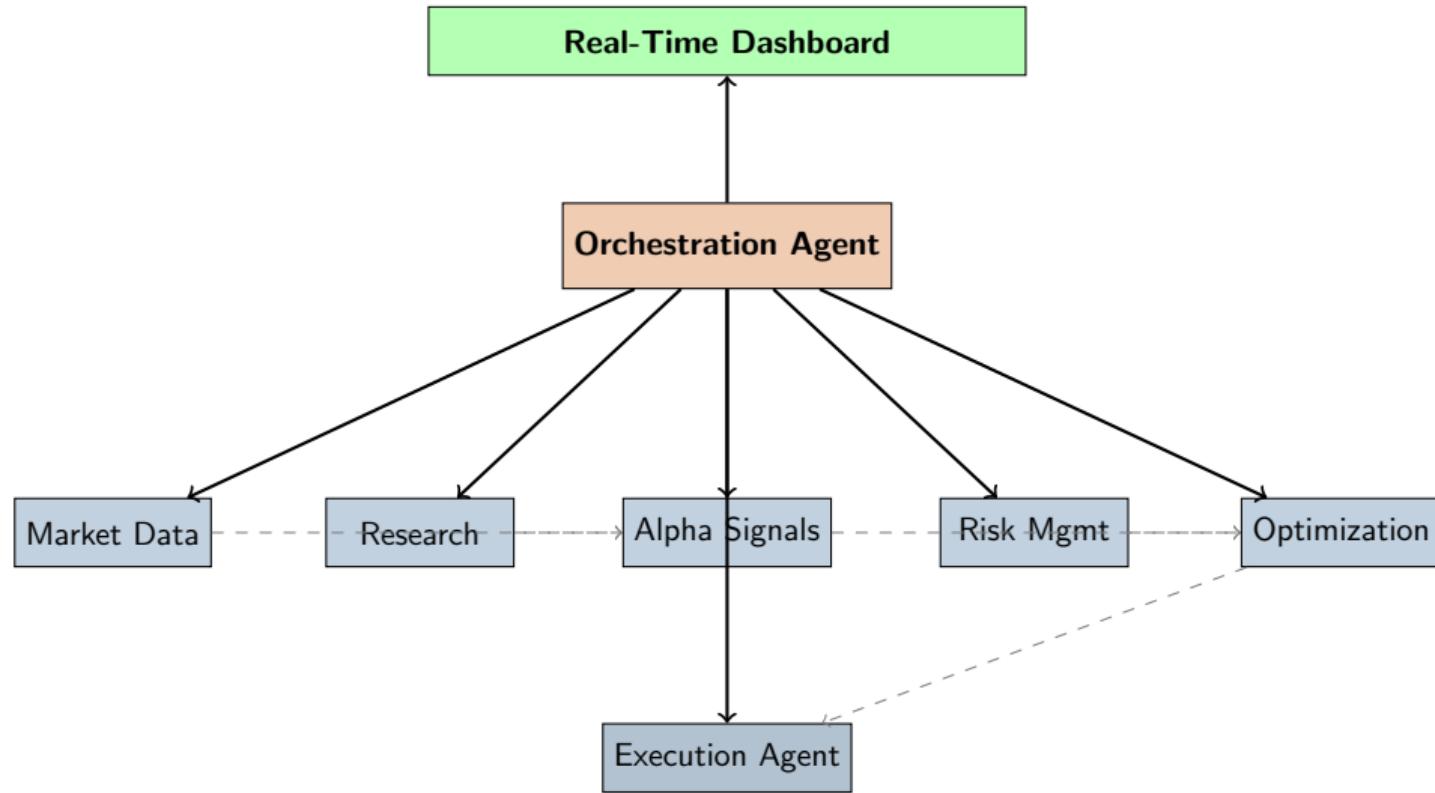
Information Flow & Decision Process

Continuous Decision Cycle

- ① Market Data Agent feeds fresh data to all agents
- ② Research Intelligence Agent provides context & sentiment
- ③ Alpha Signal Agent generates trading recommendations
- ④ Risk Management Agent evaluates against risk limits
- ⑤ Portfolio Optimization Agent creates optimal allocation
- ⑥ Orchestration Agent coordinates consensus
- ⑦ Execution Agent implements approved trades

All activities visible in real-time through the Dashboard

System Architecture Diagram



Why Agent Integration is Powerful

The Power of Collaboration

Individual agents are specialists, but **integration creates intelligence**

1. Holistic Decision Making

- Single agent: "This stock looks good based on technical analysis"
- **Integrated system:** "Technical signals are bullish BUT news sentiment is negative AND risk limits are near threshold → reduce position size by 50%"

2. Cross-Validation & Error Reduction

- Each agent provides independent perspective
- Conflicting views trigger deeper analysis
- Reduces false signals through consensus

Real-World Example: Agent Integration in Action

Scenario: Tech Stock Decision

Situation: Apple stock shows strong technical breakout signal

Without Integration (Single Agent):

- Alpha Signal Agent: "Buy 1000 shares!" → Potentially risky

With Integration (Multi-Agent):

- ① **Alpha Agent:** "Strong bullish signal, confidence 85%"
- ② **Research Agent:** "News sentiment neutral, earnings next week"
- ③ **Risk Agent:** "Tech sector exposure at 40% (limit: 45%), VaR rising"
- ④ **Optimization Agent:** "Optimal position: 500 shares given constraints"
- ⑤ **Orchestrator:** "Consensus: Buy 500 shares, set stop-loss, review after earnings"

Result: Smarter, risk-adjusted decision with context

Key Benefits of Agent Integration

1. Robustness

- If one agent fails, others continue
- Diversified decision sources
- Reduced single-point failures

2. Adaptability

- Agents learn from each other
- Market regime changes detected collectively
- Dynamic strategy adjustment

3. Interpretability

- See WHY each decision was made
- Track agent voting patterns

4. Specialization

- Each agent excels at its task
- Better than one generalist
- Easier to improve individually

5. Risk Management

- Multi-layered risk checks
- Risk agent has veto power
- Prevents catastrophic losses

6. Scalability

- Add new agents easily
- Remove underperforming ones
- Modular expansion

Information Sharing Creates Emergent Intelligence

Example Synergies

- **Data + Research:** Market data shows price spike → Research agent finds breaking news explaining it
- **Alpha + Risk:** Alpha suggests aggressive trade → Risk agent moderates based on portfolio correlation
- **Research + Optimization:** Negative news on one stock → Optimizer shifts to correlated alternatives
- **All agents + Memory:** Similar pattern failed last month → System learns and adjusts confidence

The Orchestrator's Role

Technology Stack

Backend & Core

- Python 3.10+
- FastAPI
- LangChain/CrewAI
- PostgreSQL
- Redis

Data & Analytics

- Pandas, NumPy
- yfinance, Alpha Vantage
- TA-Lib

ML & Optimization

- PyTorch/TensorFlow
- Hugging Face Transformers
- PyPortfolioOpt
- CVXPY
- QuantLib

Visualization

- Streamlit/Gradio
- Plotly
- React.js

16-Week Development Plan

Month 1: Foundation (Weeks 1-4)

- Week 1: System architecture & agent protocols
- Week 2: Market Data Agent with real-time feeds
- Week 3: Risk Management Agent (RiskIQ foundation)
- Week 4: Agent orchestration framework

Month 2: Core Intelligence (Weeks 5-8)

- Week 5: Alpha Signal Agent
- Week 6: Portfolio Optimization Agent
- Week 7: Multi-agent coordination & consensus
- Week 8: Memory & learning systems

16-Week Development Plan (continued)

Month 3: Advanced Features (Weeks 9-12)

- Week 9: AI Research Assistant with NLP
- Week 10: Execution Agent
- Week 11: Real-Time Trading Dashboard
- Week 12: Backtesting framework

Month 4: Integration & Testing (Weeks 13-16)

- Week 13: Full system integration & testing
- Week 14: Live market simulation
- Week 15: Comparative analysis & benchmarking
- Week 16: Final documentation & presentation

Week 1 Deliverable (Today's Focus)

This Week's Goal

Design system architecture and agent communication protocols

Specific Tasks:

- ① Define agent roles and responsibilities
- ② Design message passing protocol
- ③ Create system architecture diagram
- ④ Set up development environment
- ⑤ Initialize code repository
- ⑥ Define data schemas and APIs

Next Week: Build Market Data Agent with live feeds from yfinance

Performance Targets

Quantitative Goals

- **Sharpe Ratio:** > 1.5 (vs 1.0 for S&P 500)
- **Maximum Drawdown:** $< 15\%$ (vs 30-40% benchmarks)
- **Annual Returns:** 12-18% with controlled volatility
- **Win Rate:** $> 55\%$ on individual signals

System Benefits

- Reduced emotional bias through automation
- 24/7 market monitoring and risk management
- Scalable from \$10K to \$10M+ portfolios
- Consistent strategy execution

What Makes Our System Unique?

① True Multi-Agent Intelligence

- Agents reason independently & collaborate
- Like a team of portfolio managers

② Adaptive Learning

- System learns from every decision
- Strategies evolve with markets

③ Complete Transparency

- Every decision is logged and explainable
- Dashboard shows agent reasoning

④ Modular Architecture

- Easy to update individual components
- Each app works independently or integrated

Research Contributions

Novel Aspects

- **First comprehensive** multi-agent framework for portfolio optimization
- **Hybrid intelligence** combining quant models & AI analysis
- **Real-time learning** with continuous adaptation
- **Practical implementation** demonstrating live trading viability

Beyond Traditional Approaches

Our system goes beyond existing portfolio management tools by providing:

- Interpretable AI decisions
- Dynamic strategy adaptation
- Seamless integration of diverse data sources
- End-to-end automation with human oversight

Final Deliverables

① Five Working Applications

- Hosted on HuggingFace Spaces
- Fully functional and documented

② Integrated Platform

- All apps working as unified system

③ Backtesting Results

- 5+ years of historical performance

④ Source Code & Documentation

- GitHub repository with full documentation
- Architecture diagrams, API reference

⑤ Research Paper

- Academic paper documenting methodology

Summary

What We're Building

An intelligent, modular, and transparent multi-agent system for automated portfolio management that combines:

- Advanced risk analysis
- Real-time monitoring
- AI-powered research
- Portfolio optimization
- Coordinated agent decision-making

Impact

This project demonstrates how modern AI and multi-agent systems can be applied to quantitative finance to build robust, interpretable, and scalable solutions

Next Steps

Immediate Actions (Week 1):

- ① Finalize architecture design
- ② Set up development environment
- ③ Create agent communication protocol
- ④ Initialize GitHub repository
- ⑤ Begin Market Data Agent development

Next Presentation (Week 2):

- Demo of Market Data Agent
- Real-time data feeds working
- Data validation & preprocessing

Thank You!

Questions?

Ashwini (MA25M006)
Dibyendu Sarkar (MA25M011)
Jyoti Ranjan Sethi (MA25M013)

Backup: Risk Metrics

Value at Risk (VaR)

Maximum expected loss over a time period at a given confidence level

$$\text{VaR}_\alpha = \inf\{I \in \mathbb{R} : P(L > I) \leq 1 - \alpha\}$$

Conditional VaR (CVaR)

Expected loss given that loss exceeds VaR

$$\text{CVaR}_\alpha = E[L | L > \text{VaR}_\alpha]$$

Sharpe Ratio

Risk-adjusted return metric

$$\text{Sharpe} = \frac{E[R_p - R_f]}{\sigma_p}$$

Backup: Portfolio Optimization

Mean-Variance Optimization

$$\begin{aligned} \min_w \quad & w^T \Sigma w \\ \text{s.t.} \quad & w^T \mu \geq r_{\min}, \quad \sum w_i = 1, \quad w_i \geq 0 \end{aligned}$$

Sharpe Ratio Maximization

$$\max_w \quad \frac{w^T \mu - r_f}{\sqrt{w^T \Sigma w}}$$

Where:

- w = portfolio weights
- μ = expected returns
- Σ = covariance matrix
- r_f = risk-free rate

Backup: Agent Communication Protocol

Message Structure

```
{  
  "agent_id": "risk_agent",  
  "timestamp": "2025-02-12T10:30:00",  
  "message_type": "risk_alert",  
  "content": {  
    "metric": "VaR",  
    "value": 0.05,  
    "threshold": 0.03,  
    "action": "reduce_exposure"  
  }  
}
```